

# Scalable Data Engineering, Exercise 2

HENRY HAUSTEIN

## Task 1

- (a) False, for some reporting functions we need `OVER`
- (b) True
- (c) False, this does `CUBE`
- (d) True
- (e) True
- (f) False, SQL integer division returns an integer
- (g) False, `ORDER BY` in an `OVER` clause has no effect on the sorting of the result. And an `ORDER BY` outside an `OVER` only sorts the result.
- (h) False, MDX operates on data cubes. This can sometimes work on relational databases but it's not always the case.

## Task 2

- (a) SQL:

```
1 SELECT avg(ORDERS.O_TOTALPRICE) FROM ORDERS
```

	avg(ORDERS.O_TOTALPRICE)
1	151219.537631631

- (b) SQL (I'm using Sqlite for this task, it doesn't support `ROLLUP`):

```
1 SELECT REGION.R_NAME ,
2 FROM REGION , ORDERS , CUSTOMER , NATION
3 WHERE
4     ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY AND
5     CUSTOMER.C_NATIONKEY = NATION.N_NATIONKEY AND
6     NATION.N_REGIONKEY = REGION.R_REGIONKEY
7 GROUP BY REGION.R_NAME
8 UNION
9 SELECT NULL , sum(ORDERS.O_TOTALPRICE) FROM ORDERS
```

	<b>R_NAME</b>	<b>sum(ORDERS.O_TOTALPRICE)</b>
<b>1</b>	NULL	226829306447.447
<b>2</b>	AFRICA	45230223902.2213
<b>3</b>	AMERICA	45306943255.2101
<b>4</b>	ASIA	45613415042.5609
<b>5</b>	EUROPE	45793265459.7115
<b>6</b>	MIDDLE EAST	44885458787.7611